



University of Pisa

Department of Information Engineering

Master's Degree Artificial Intelligence and Data Engineering

Cloud Computing Project: Inverted Index

Professors:

Carlo Vallati

Carlo Puliafito

Students:

Andrea Zanin

Lorenzo Valtriani

Samuele Marchi

ACADEMIC YEAR 2024/2025

Contents

1	Introduction	2
1.1	Problem presentation	2
1.2	Dataset description	2
2	Pseudocode	3
2.1	External combiner version	3
2.2	In-Mapper combining version	4
3	Performance Evaluation	5
3.1	Different input split sizes	5
3.2	External combiner vs In-Mapper combining	6
3.3	Different number of reducers	6
3.4	Time comparison between different solutions	7
3.5	Resources comparison between different solutions	8

Introduction

1.1 Problem presentation

This project, developed within the field of *cloud computing*, focuses on implementing an *inverted index* on a large dataset composed by small-sized scientific papers (on the order of kilobytes).

The application was implemented in three ways: using **Hadoop MapReduce**, with both a *combiner* and an *in-mapper combining* variant; a distributed version with **Apache Spark** and a non parallel version developed in **Python**. The output maps each word (**word**) to the list of files in which it appears, along with the number of occurrences per file, in the following format:

```
cloud          file1.txt:3  file2.txt:4
computing      file2.txt:2  file3.txt:1
```

The project concludes with a comparative analysis of the three approaches, focusing on *execution time* and *resource management*.

1.2 Dataset description

The dataset was collected using a custom Python web scraping script that automatically extracted scientific articles from the [arXiv website](#). The original PDF files were converted to plain text to enable processing with distributed frameworks such as MapReduce and Apache Spark.

To assess the performance and scalability of the implementations, six dataset versions were generated, varying in total size and number of files.

Dataset name	Size	Number of files
dataset_kilobytes	276 KB	6
dataset_megabytes	94 MB	1,622
dataset_megabytes2	353 MB	5,609
dataset_megabytes3	706 MB	11,218
dataset_gigabytes	1.1 GB	19,064
dataset_gigabytes2	2.2 GB	57,192

This segmentation allows for evaluating algorithm efficiency under increasing workload conditions, similar to those in modern search engines and large-scale text processing systems.

Pseudocode

2.1 External combiner version

Algorithm 1 Mapper

```
1: class INVERTEDINDEXMAPPER EXTENDS MAPPER
2:   Variables: stopWords  $\leftarrow$  new empty list
3:   method SETUP(context)
4:     stopWordsFile  $\leftarrow$  load stopwords.txt
5:     for all line in stopWordsFile do
6:       w  $\leftarrow$  cleaned lowercase line
7:       if w  $\neq$  empty then
8:         add w to stopWords
9:       end if
10:    end for
11:  end method
12:
13:  method MAP(key, value, context)
14:    fName  $\leftarrow$  get file name from key
15:    tokens  $\leftarrow$  Preprocess(value)
16:    for all t in tokens do
17:      EMIT(t, (fName, 1))
18:    end for
19:  end method
20: end class
```

Algorithm 2 Combiner

```
1: class INVERTEDINDEXCOMBINER EXTENDS REDUCER
2:   method REDUCE(key, values, context)
3:     fileCounts  $\leftarrow$  new AssociativeArray
4:     for all val in values do
5:       fileName  $\leftarrow$  get the file name from val
6:       count  $\leftarrow$  get the counter from val
7:       fileCounts{fileName}  $\leftarrow$  fileCounts{fileName} + count
8:     end for
9:     for all file in fileCounts do
10:      count  $\leftarrow$  fileCounts{file}
11:      EMIT(key, (file, fileCounts{file}))
12:    end for
13:  end method
14: end class
```

Algorithm 3 Reducer

```
1: class INVERTEDINDEXREDUCER EXTENDS REDUCER
2:   method REDUCE(key, values, context)
3:     fileCounts  $\leftarrow$  new AssociativeArray
4:     result  $\leftarrow$  new text
5:     for all val in values do
6:       fileName  $\leftarrow$  get the file name from val
7:       count  $\leftarrow$  get the counter from val
8:       fileCounts{fileName}  $\leftarrow$  fileCounts{fileName} + count
9:     end for
10:    for all file in fileCounts do
11:      result  $\leftarrow$  concat(file:fileCounts{file})
12:    end for
13:    EMIT(key, result)
14:  end method
15: end class
```

2.2 In-Mapper combining version

Algorithm 4 Mapper with In-Mapper Combining

```

1: class INVERTEDINDEXMAPPER EXTENDS MAPPER
2:   Variables: stopWords  $\leftarrow$  empty list, wordMap  $\leftarrow$  empty map, fileName  $\leftarrow$  empty
3:   method SETUP(context)
4:     stopWordsFile  $\leftarrow$  load stopwords.txt
5:     for all line in stopWordsFile do
6:       w  $\leftarrow$  cleaned lowercase line
7:       if w  $\neq$  empty then
8:         add w to stopWords
9:       end if
10:    end for
11:  end method
12:
13:  method MAP(key, value, context)
14:    curFile  $\leftarrow$  get file name from key
15:    if fileName is null then
16:      fileName  $\leftarrow$  curFile
17:    end if
18:    if fileName  $\neq$  curFile then
19:      Flush(context), fileName  $\leftarrow$  curFile
20:    end if
21:    for all t in Preprocess(value) do
22:      if t  $\neq$  empty then
23:        wordMap{t}  $\leftarrow$  wordMap{t} + 1
24:      end if
25:    end for
26:  end method
27:
28:  method CLEANUP(context)
29:    Flush(context)
30:  end method
31:
32:  method FLUSH(context)
33:    for all w in wordMap do
34:      EMIT(w, fileName:wordMap{w})
35:    end for
36:    clear wordMap
37:  end method
38: end class

```

Algorithm 5 Reducer

```

1: class INVERTEDINDEXREDUCER EXTENDS REDUCER
2:   method REDUCE(key, values, context)
3:     result  $\leftarrow$  convert values in serializable form
4:     EMIT(key, result)
5:   end method
6: end class

```

Performance Evaluation

Before analyzing the performance of the different applications, it is important to note that the cluster, provided by the University of Pisa, consists of three VMs, each equipped with **7 GB of RAM** and **40 GB of disk**. Each VM has been configured with the following Hadoop and Spark settings:

Configuration	Value
yarn.nodemanager.resource.memory-mb	1536
yarn.scheduler.maximum-allocation-mb	1536

Configuration	Value
yarn.app.mapreduce.am.resource.mb	512
mapreduce.map.memory.mb	256
mapreduce.reduce.memory.mb	256


Configuration	Value
spark.master	yarn
spark.driver.memory	512m
spark.yarn.am.memory	512m
spark.executor.memory	512m

3.1 Different input split sizes

In this section, the application was experimentally evaluated by varying the *input split size*, thereby affecting the number of mappers. The use of `CombineTextInputFormat` allowed multiple small files to be grouped into a single input split, optimizing processing in scenarios with large numbers of small files.

Smaller input splits result in *fewer spilled records per map task*, reducing disk I/O and *lowering the average mapper execution time*. However, this setup leads to a *longer overall job duration* and a *noticeable increase in memory usage*.

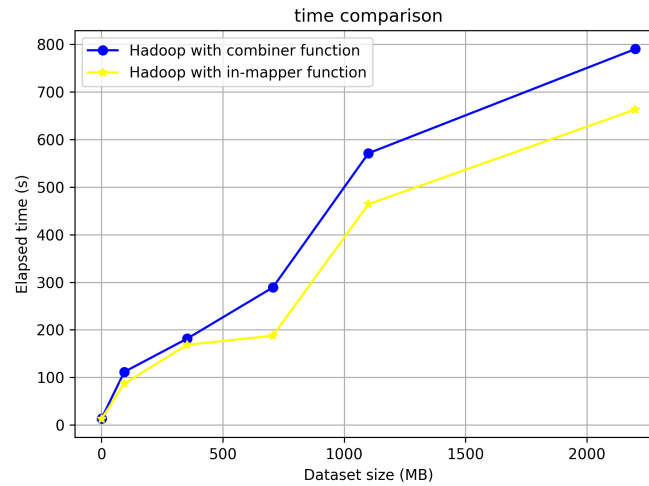
In conclusion, for applications dealing with many small files, as in this case, using `CombineTextInputFormat` with an input split size equal to the HDFS block size is an effective strategy.



3.2 External combiner vs In-Mapper combining

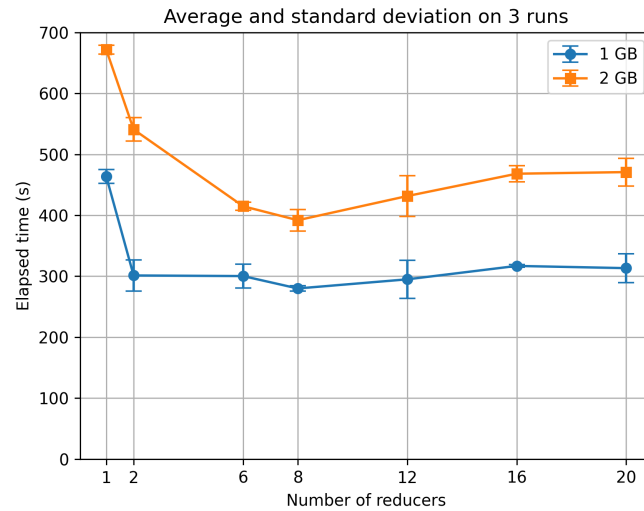
We ran the MapReduce application, first using an external combiner and then implementing in-mapper combining. We then compared the performance of the two approaches.

The In-Mapper solution offers several advantages. In all the cases analyzed, its execution time decreases, both in terms of CPU time and overall elapsed time. Additionally, memory utilization is reduced, not only in terms of peak memory usage but also in terms of the total memory consumed.



3.3 Different number of reducers


We now analyze how performance changes as the number of reducers increases. Specifically, we first tested the application on the 1 GB dataset and then on the 2 GB dataset, increasing the number of reducers up to 20 and performing 3 different runs for each configuration. This approach allowed us to compute the average elapsed time and its standard deviation for each setting. The results show that, for both datasets, **the fastest execution was achieved when using 8 reducers** which seems to be a good compromise *between load balancing and overhead*.



3.4 Time comparison between different solutions

Regarding execution time, it can be observed that **Spark proves to be the fastest solution**. This is particularly true when dividing the RDDs into partitions of approximately 750 files. This number was determined experimentally, aiming to reduce the number of partitions to lower overhead, while still leveraging only RAM for data processing, thus taking advantage of Spark's architecture.

It is evident that, although the non-parallel solution performs very well on small datasets, its execution time increases significantly as the data size grows. In contrast, both **Hadoop** and **Spark** exhibit a growth trend that is similar to a logarithmic curve, making them substantially more efficient when handling large-scale datasets.



3.5 Resources comparison between different solutions

Regarding the number of resources allocated by each of the compared solutions, the analysis considers the aggregate resource allocation, a metric available for both Hadoop and Spark. This value represents the **total amount of memory allocated over time**, calculated as the sum of the memory used by each container multiplied by the duration of its activity. For small datasets, resource usage remains nearly identical across all approaches. However, with increasing dataset size, a clear trend emerges: **the Hadoop solution with the external combiner allocates a significantly higher number of resources** compared to the other two solutions.

